

Mid-PhD Defense

Paul Dubois

TheraPanacea
MICS, CentraleSupélec
Institut du Cancer de Montpellier

21st June 2023

Outline

Introduction

Cancer treatments

Radiotherapy

Multi-Leaf Collimator

V-MAT Scheme

IMRT Scheme

Step-and-Shoot

Sliding-windows

Radiotherapy Workflow

Problem Statement

Optimization workflow

Fluence discretization

FMO problem

Formulation

Optimization

Early results

Optimizers Review

Meta-Optimization

Dose Distances

Dose Clustering

Future work

Reinforcement Learning

General Setup

Planned Network Architecture

Trick

Challenges

Others

Teaching

Doctoral Training

References

Cancer treatments

Surgery



- +: Safe (little damage to healthy tissues)
- : Tumor needs to be localized & accessible

Chemotherapy



- : Heavy medicine on all the body
- +: Tumor does **not** need to be localized

Cancer treatments

Surgery



+: Safe

-: Tumor needs to be localized

Radiotherapy



+: Relatively safe (most tissues are spared)

-: Tumor needs to be (relatively) localized

Chemotherapy



Medicine on all the body

does not need to be localized

Multi-Leaf Collimator



V-MAT Irradiation Technique

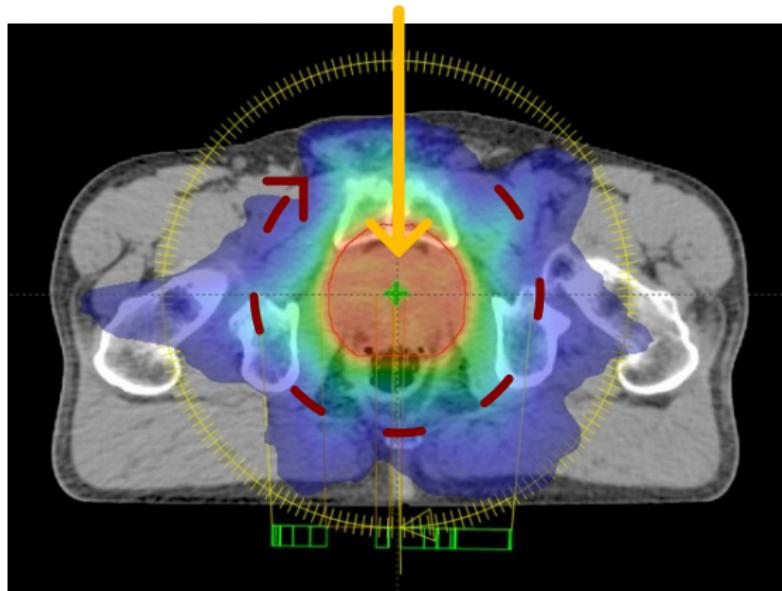


Figure: Typical V-Mat dose slice.

IMRT Irradiation Technique

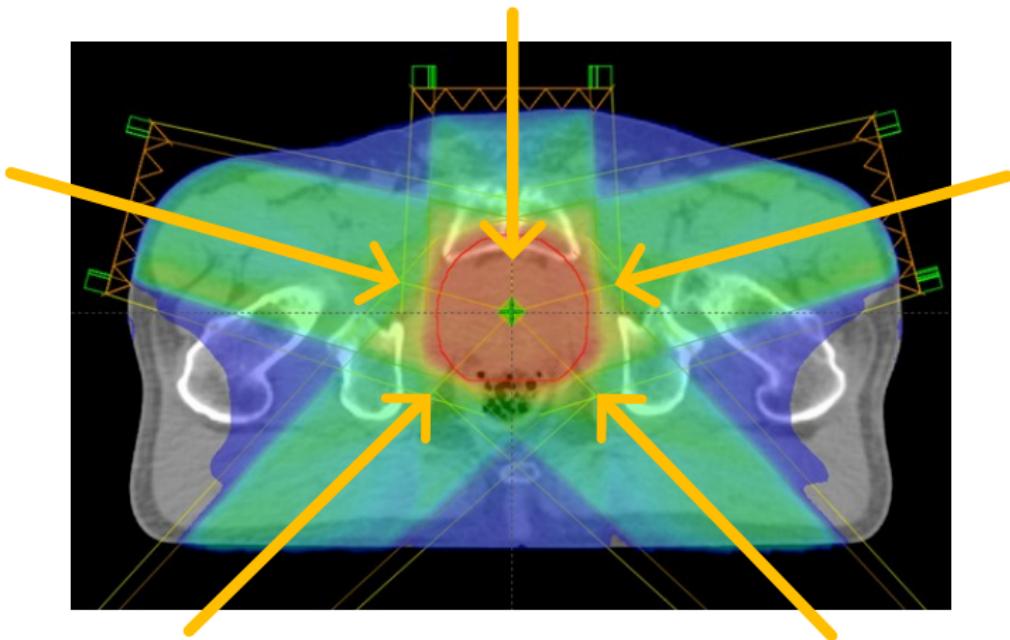


Figure: Typical 5 beams IMRT dose slice.

Step-and-Shoot (1/3)

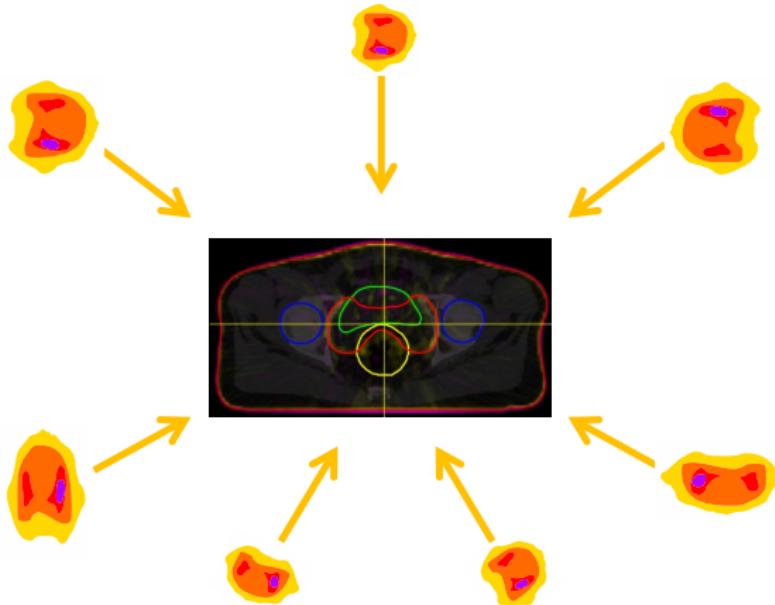


Figure: Optimal Continuous Fluence.

Step-and-Shoot (2/3)

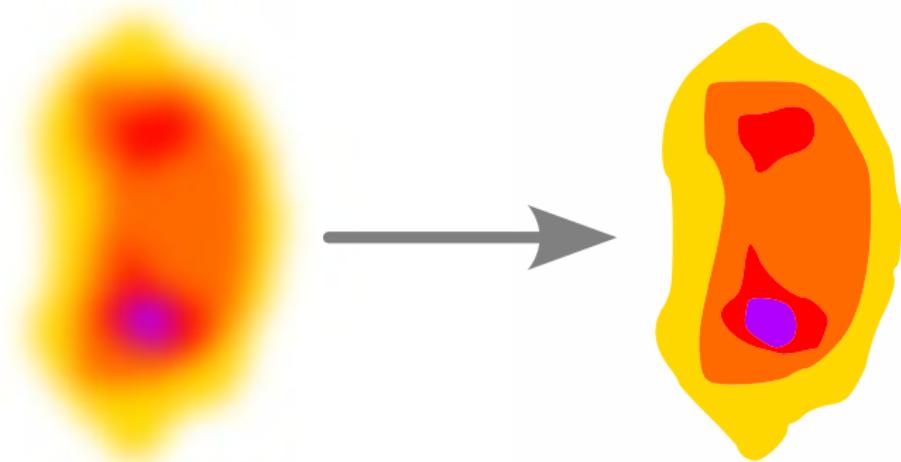


Figure: Discretizing the Fluence.

Step-and-Shoot (3/3)

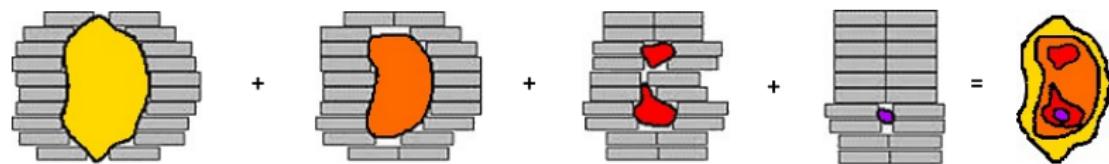


Figure: Delivering Discrete Fluence.

Sliding-Windows (1/3)



Figure: Continuous Fluence to Bixel Fluence.

Sliding-Windows (2/3)

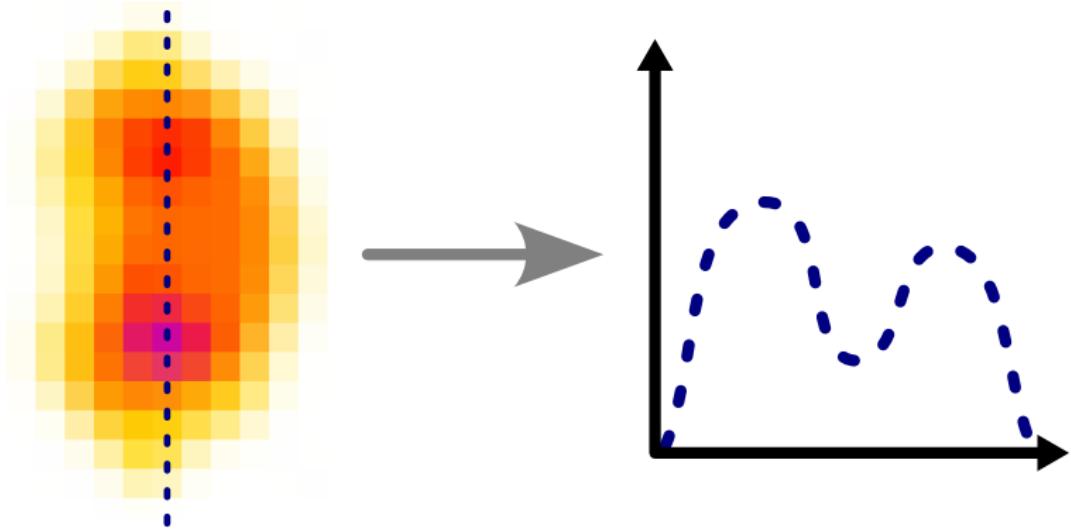


Figure: Bixel Fluence to Row/Column Curves.

Sliding-Windows (3/3)

Convert rows/columns fluence curves to leafs motions.



(<https://mics-lab.github.io/PresentationJuin2023MICS/demo>)

Radiotherapy Workflow



Radiotherapy Workflow



Automatic Dose Optimization for Radiotherapy



Problem Formulation

IMRT

Bixel values:

$$x_{i,j}^{\theta} \geq 0, \text{ for } \theta \in \Theta \text{ and } 1 \leq i,j \leq 20^1$$

usually concatenated to a single bixels-value vector x .

Dose calculation:

$$\mathbf{y} = L\mathbf{x} \text{ with } L \text{ (pre-calculated) dose-influence (DI) matrix}$$

¹20x20 is a typical bixel discretization

Problem Formulation

IMRT (bis)

Objective for *maximum* constraint c on structure s , dose d :

$$f_c(\mathbf{y}) = \frac{1}{|\mathcal{V}|} \sum_{v \in \mathcal{V}} (\mathbf{y}_v - d)_+^2$$

(reverse sign for minimal constraint).

Final objective:

$$f(\mathbf{y}) = \sum_{c \in \mathcal{C}} w_c f_c(\mathbf{y})$$

with w_c the weight of constraint c .

Problem Optimization

Optimizer review



Figure: Typical prostate case.

<https://arxiv.org/abs/2305.18014>

Problem Optimization

Optimizer review (bis)



Figure: Typical prostate case.

<https://arxiv.org/abs/2305.18014>

Meta-Optimization

Usual optimization

$$\min_{\mathbf{x}} f(\mathbf{x}, w) \text{ s.t. } \mathbf{x} > 0$$

... and fine-tune w until the dose is clinically acceptable.

Meta optimization

$$\min_w \left\{ \min_{\mathbf{x}} f(\mathbf{x}, w) \text{ s.t. } \mathbf{x} > 0 \right\}$$

... still need to fine-tune the parameters (learning rate, momentum, etc...) of the meta-optimizer.

Dose Distances

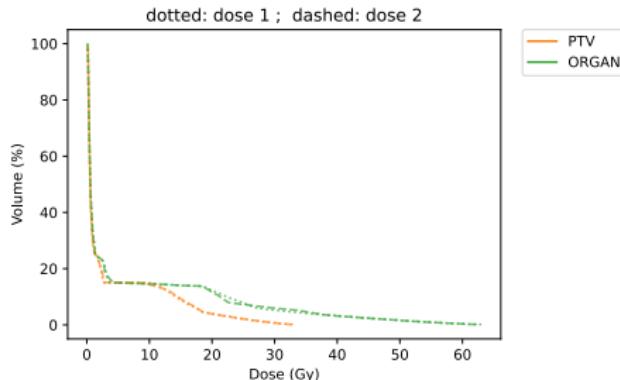
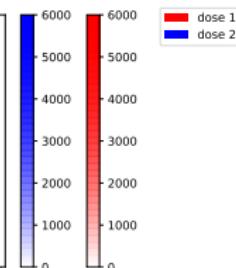
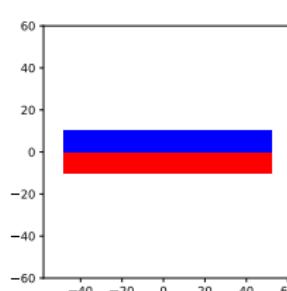
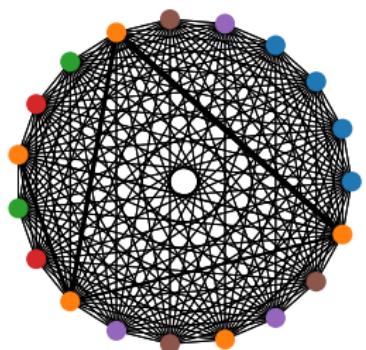


Figure: Example of two doses that have the same clinical effect (measured from the DVHs), but very different voxel-wise dose values.

Dose Clustering



(a) (Circular Layout)



(b) (Spring Layout)

Figure: Doses Network

edges width \propto edge weight $\propto 1/\text{distance}$

node's color reflects community attribution

Dose Clustering

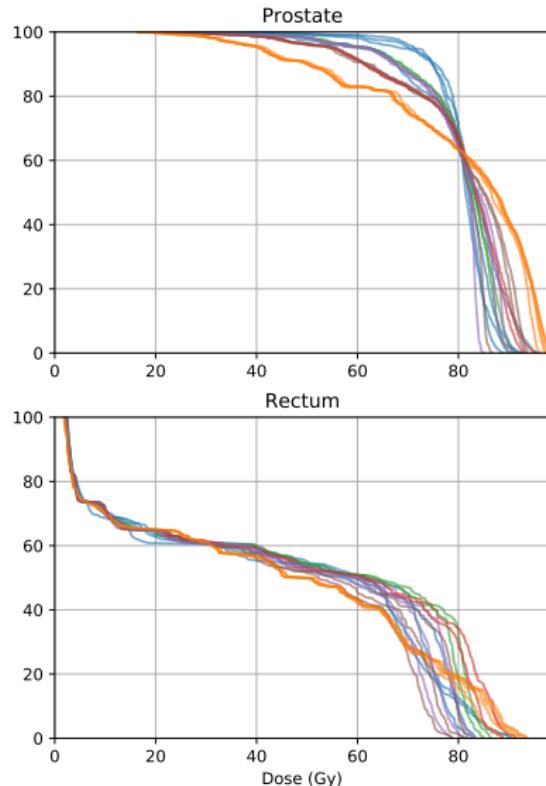


Figure: Dose-Volume Histogram

Dose Clustering

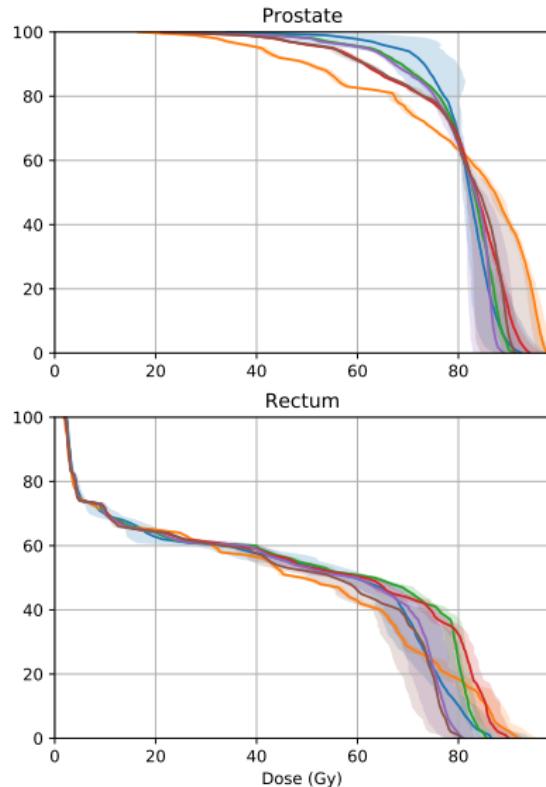
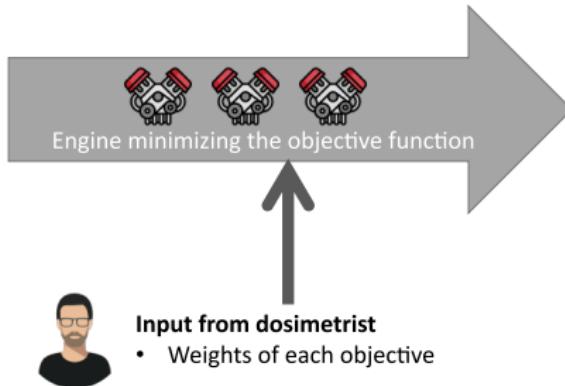


Figure: Dose-Volume Histogram Standard Deviation per Community

Current Workflow

Patient data

- CT scan
- OARs & PTVs contours
- doctors' objectives



Treatment plan

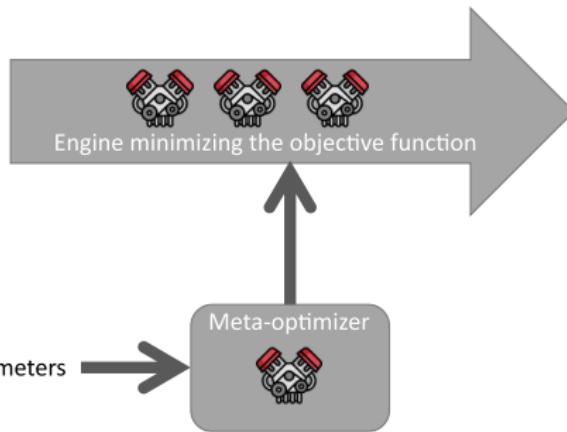
- Fluences
- Leaf movements
- Dose per organ

1st Automatization try “baseline”

Patient data

- CT scan
- OARs & PTVs contours
- doctors' objectives

Fine tuning of hyper parameters



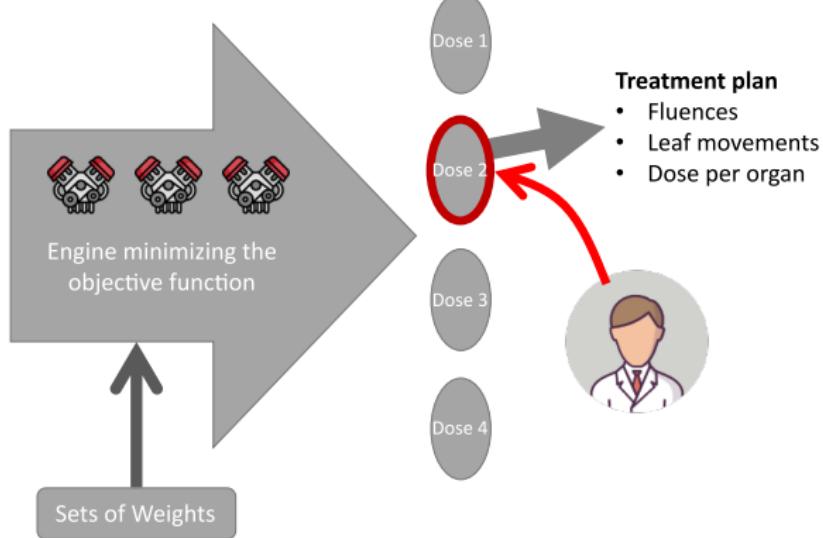
Treatment plan

- Fluences
- Leaf movements
- Dose per organ

1st (bis) Automatization try

Patient data

- CT scan
- OARs & PTVs contours
- doctors' objectives

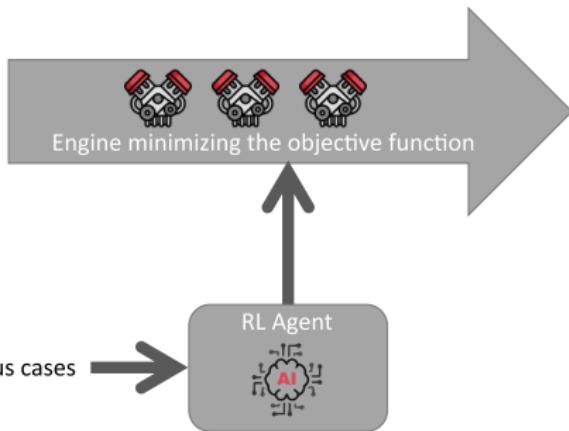


2nd Automatization try (*work in progress*)

Patient data

- CT scan
- OARs & PTVs contours
- doctors' objectives

Learning based on previous cases



Treatment plan

- Fluences
- Leaf movements
- Dose per organ

Reinforcement Learning Setup

Agent A network.

Env. The current dose/weights, the CT scan, structures contours.

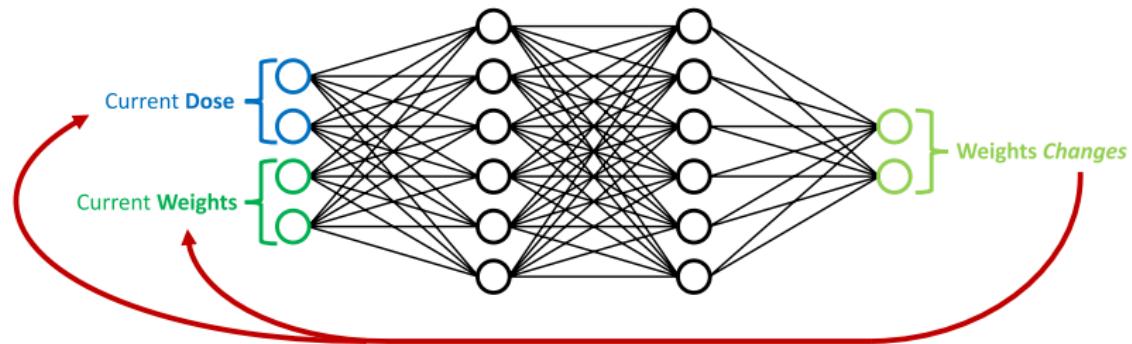
State The agent will only access the the current dose/weights.

Action Changing the set of weights.

Reward The (DVH) distance between the current dose and the one that was actually used.

Policy The value changes of the sets of weights.

Planned Network Architecture



Trick: encode the dose to smaller space

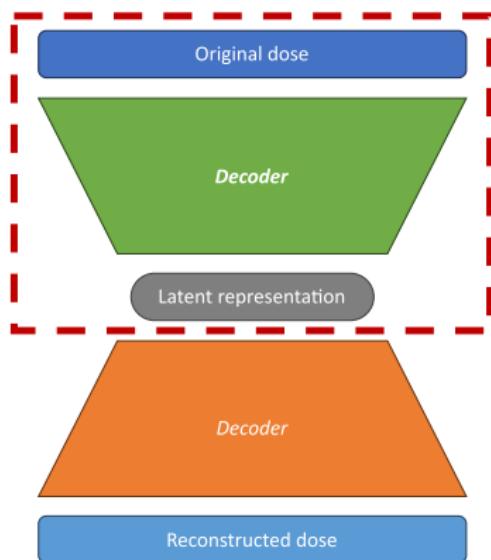


Figure: Dose Auto-Encoder Architecture

Challenges

- ▶ normalize the body scan size
- ▶ normalize the structures per anatomy
- ▶ normalize the constraints per anatomy

Teaching

Lectures:

- ▶ Mathematics Refresher Course for DSBA (M2 students) 2021
- ▶ Deep Learning for HSB (3rd year students) 2023

TDs:

- ▶ Coding Weeks (1st year) 2021
- ▶ Optimization (1st year) 2021
- ▶ Visual recognition (3rd year) 2022
- ▶ Coding Weeks (1st year) 2022
- ▶ Algorithm and Complexity (2nd year) 2022/2023

Doctoral Training

- ▶ ED INTERFACES - Journée de Rentrée 2022 (12 janvier 2023)
- ▶ Math On Mars (06 mai 2022) Info@lèze
- ▶ Asymmetric Cryptography (23 septembre 2022) Info@lèze
- ▶ Genetic Algorithms (10 juin 2022) Info@lèze
- ▶ Math With Jupyter (18 décembre 2021) Info@lèze
- ▶ Writing skills in Science ADVANCED [Eng] (10 mai 2022)
- ▶ AI 4 Health (10 janvier 2022 - 14 janvier 2022)

Total participation: **109/125** heures; 7 modules

Total des Crédits/Points de Thèse: **22/25**

References I

-  Paul Dubois.
Radiotherapy dosimetry: A review on open-source optimizer, 2023.
-  Gary A. Ezzell, James M. Galvin, Daniel Low, Jatinder R. Palta, Isaac Rosen, Michael B. Sharpe, Ping Xia, Ying Xiao, Lei Xing, and Cedric X. Yu.
Guidance document on delivery, treatment planning, and clinical implementation of imrt: Report of the imrt subcommittee of the aapm radiation therapy committee.
Medical Physics, 30(8):2089–2115, 2003.
-  Thomas P. Kole, Osarhieme Aghayere, Jason Kwah, Ellen D. Yorke, and Karyn A. Goodman.
Comparison of heart and coronary artery doses associated with intensity-modulated radiotherapy versus three-dimensional conformal radiotherapy for distal esophageal cancer.

References II

*International Journal of Radiation Oncology*Biology*Physics*, 83(5):1580–1586, 2012.

-  Dinesh Kumar Mynampati, Ravindra Yaparpalvi, Linda Hong, Hsiang-Chi Kuo, and Dennis Mah.
Application of AAPM TG 119 to volumetric arc therapy (VMAT).
Journal of Applied Clinical Medical Physics, 13(5):108–116, 2012.
-  David Palma, Emily Vollans, Kerry James, Sandy Nakano, Vitali Moiseenko, Richard Shaffer, Michael McKenzie, James Morris, and Karl Otto.
Volumetric modulated arc therapy for delivery of prostate radiotherapy: Comparison with intensity-modulated radiotherapy and three-dimensional conformal radiotherapy.
*International Journal of Radiation Oncology*Biology*Physics*, 72(4):996–1001, 2008.

References III

-  Bortfeld Thomas.
The number of beams in imrt—theoretical investigations and implications for single-arc imrt.
Physics in Medicine & Biology, 55:83 – 97, 2009.
-  Lynn J. Verhey.
Comparison of three-dimensional conformal radiation therapy and intensity-modulated radiation therapy systems.
Seminars in Radiation Oncology, 9(1):78–98, 1999.
Radiation Therapy Treatment Optimization.
-  S Webb.
The physical basis of imrt and inverse planning.
The British Journal of Radiology, 76(910):678–689, 2003.
PMID: 14512327.